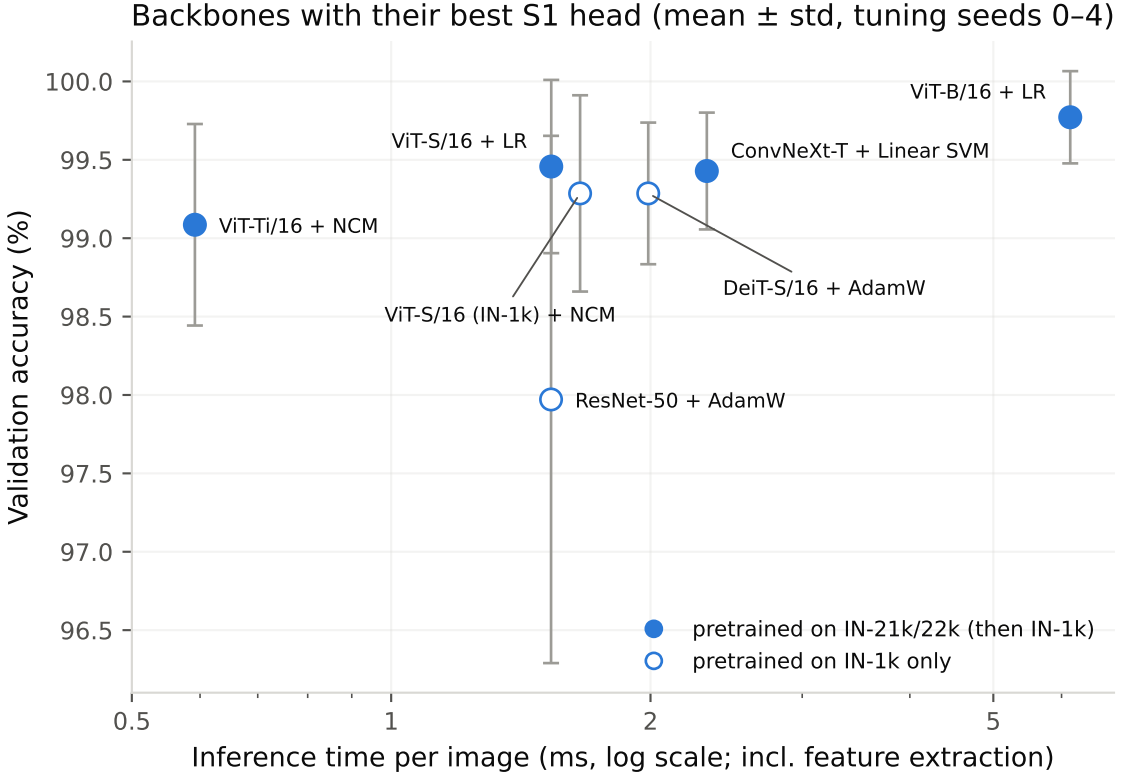
Backbones with their best S1 head (mean ± std, tuning seeds 0–4)
Validation accuracy (%)
100.0
99.5
99.0
98.5
98.0
97.5
97.0
96.5
0.5
1
2
5
Inference time per image (ms, log scale; incl. feature extraction)
ViT-Ti/16 + NCM
ViT-S/16 + LR
ViT-S/16 (IN-1k) + NCM
ResNet-50 + AdamW
ConvNeXt-T + Linear SVM
DeiT-S/16 + AdamW
ViT-B/16 + LR
pretrained on IN-21k/22k (then IN-1k)
pretrained on IN-1k only