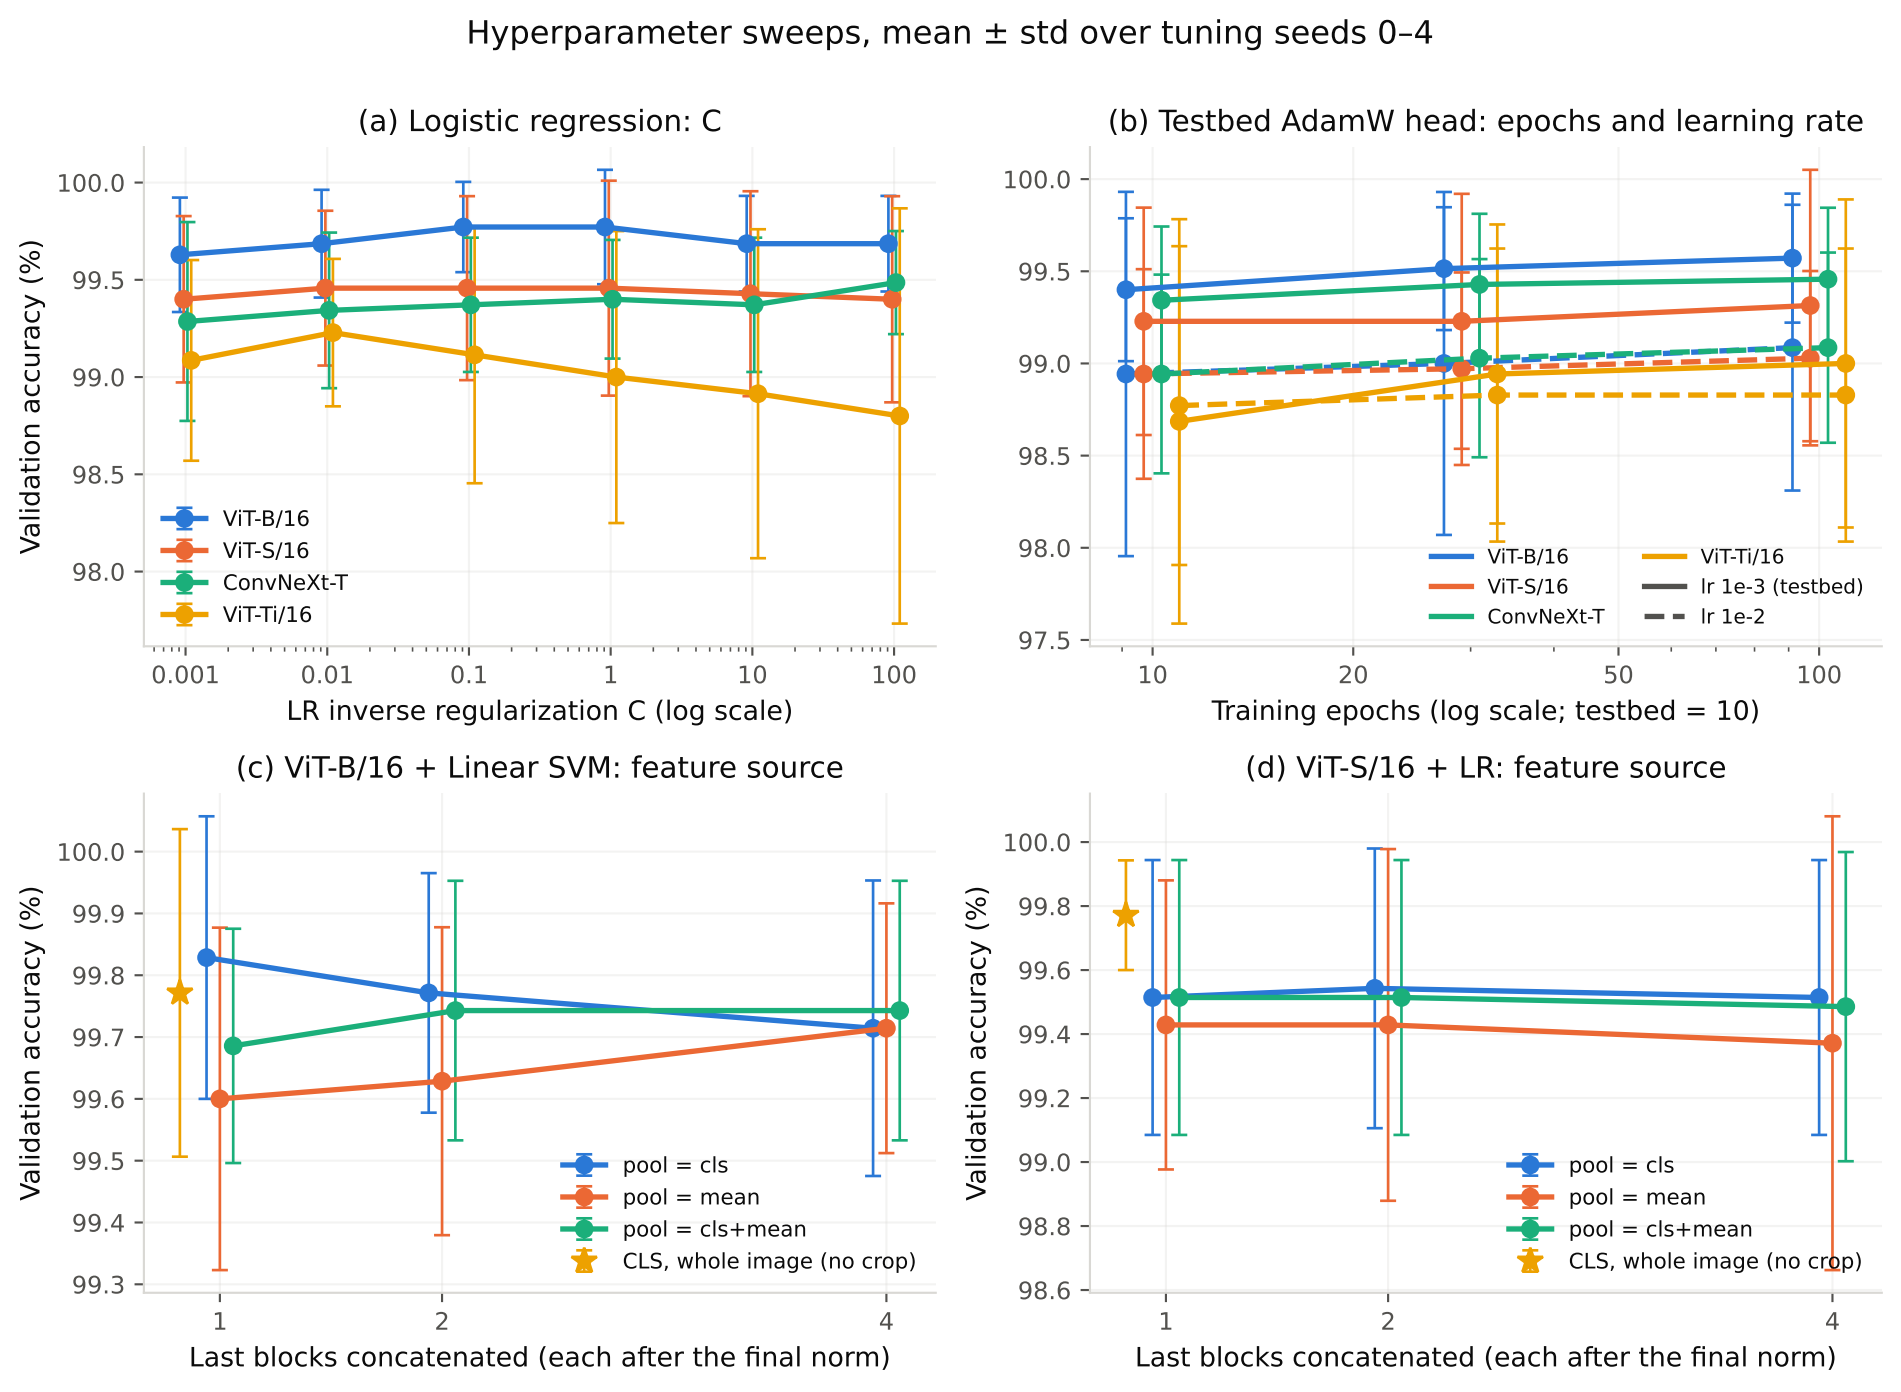
Hyperparameter sweeps, mean ± std over tuning seeds 0–4
(a) Logistic regression: C
Validation accuracy (%)
LR inverse regularization C (log scale)
ViT-B/16
ViT-S/16
ConvNeXt-T
ViT-Ti/16
(b) Testbed AdamW head: epochs and learning rate
Training epochs (log scale; testbed = 10)
ViT-B/16
ViT-S/16
ConvNeXt-T
ViT-Ti/16
lr 1e-3 (testbed)
lr 1e-2
(c) ViT-B/16 + Linear SVM: feature source
Validation accuracy (%)
Last blocks concatenated (each after the final norm)
pool = cls
pool = mean
pool = cls+mean
CLS, whole image (no crop)
(d) ViT-S/16 + LR: feature source
Validation accuracy (%)
Last blocks concatenated (each after the final norm)
pool = cls
pool = mean
pool = cls+mean
CLS, whole image (no crop)